

RETAILSPHERE (PTY) LTD

Data Platform Programme

Business Process Analysis Report

Step 2 Deliverable — Understanding Business Operations Across All Six Domains

Document ID	BPA-001
Version	1.0 — Final. All 5 interviews complete.
Date	1 May 2026
Author	Thozamile Mbalo — Data Solutions Architect
Reviewed by	Nomvula Mahlangu — Chief Data Officer
Status	Draft — Pending CDO Sign-off
Interviews	5/5 Complete — All domain process owners interviewed

1. Purpose

This document presents the findings of Step 2 — Business Process Analysis — of the RetailSphere Data Platform Programme. The purpose of Step 2 is to understand what each business domain actually does, what decisions stakeholders make, what questions they cannot currently answer, and what success looks like in 12 months. These findings form the foundation of the functional requirements in Step 3.

Five domain process owners were interviewed over the period April to May 2026. The CDO, Nomvula Mahlangu, observed all sessions. Each interview followed a structured seven-heading framework: process owner profile, key business processes, decisions made, questions they cannot answer, what good looks like in 12 months, data sources and quality issues, and architectural implications.

This report supersedes BPA-001 Draft v0.3, which contained three completed interview sections. Version 1.0 includes all five domain interviews and four formal findings.

2. Interview Status

Domain	Process Owner	Title	Date	Status
Procurement	Sipho Dlamini	Head of Procurement	17 Apr 2026	Complete
Inventory Mgmt	Zanele Mokoena	National Inventory Manager	17 Apr 2026	Complete

Domain	Process Owner	Title	Date	Status
Sales / Marketing	Rethabile Sithole	Chief Commercial Officer	17 Apr 2026	Complete
Customer Loyalty	Priya Naidoo	Head of CRM and Loyalty	17 Apr 2026	Complete
Finance	Johan van der Merwe	Group Financial Controller	17 Apr 2026	Complete

3. Domain Interview Findings

3.1 Procurement — Sipho Dlamini, Head of Procurement

Interview complete — 17 April 2026 | ADRs triggered: ADR-001, ADR-003

Process Owner Profile

Sipho Dlamini is the Head of Procurement at RetailSphere, responsible for managing the end-to-end procurement process across all 87 stores and the distribution centre, working with 340 active suppliers. He is a direct communicator who has been working around SAP's data limitations for years and walked into the interview carrying annotated SAP reports covered in red pen.

Key Business Processes

The procurement process follows six stages: demand signal via store POS or SAP replenishment report triggering a Purchase Requisition; Purchase Order creation where the buyer selects a supplier — this decision is undocumented and lives in the buyer's head; supplier confirmation via email with manual SAP update creating a data quality risk; Goods Receipt Note capture in SAP which is currently batch-captured at end of day creating timestamp unreliability; three-way invoice matching in SAP where exceptions are managed in Excel by Johan's Finance team; and payment runs on 30, 60, or 90 day supplier terms.

A critical cross-domain dependency exists — GRN data quality is owned by Zanele Mokoena's Inventory team, not Sipho. The warehouse team capturing GRNs reports to Zanele. Sipho's procurement data quality depends on a team he does not manage.

Decisions Made — Daily, Weekly, Monthly

Daily: Buyers review open POs and follow up on supplier confirmations. The exceptions queue is monitored for failed invoice matches. Weekly: Procurement savings are tracked through manual calculation. Supplier delivery performance is reviewed, though this metric is currently unavailable in SAP. Monthly: The board requires procurement savings as a percentage of total spend — a number Sipho's team currently calculates manually over three days in Excel.

Questions Sipho Cannot Currently Answer

- Which suppliers are consistently late — and by how many days on average?
- Which suppliers have quality rejection rates above acceptable threshold?
- What is the price variance between contracted rates and what was actually paid?
- What are our Days Payable Outstanding by supplier?
- What is procurement savings as a percentage of total spend — without three days of manual Excel work?

What Good Looks Like in 12 Months

Sipho's Monday morning dashboard — five metrics visible without manual calculation:

- Supplier on-time delivery rate
- Invoice accuracy rate
- Price variance against contracted rates
- Days Payable Outstanding
- Procurement savings as percentage of total spend

Data Sources and Quality Issues

The primary source system is SAP ECC 6.0, which holds purchase orders, goods receipts, invoices, payments, and supplier master data. A secondary source exists — 340 supplier contracts stored as unstructured PDF documents on a shared drive, not in SAP.

Issue 1 — GRN Timestamps Unreliable

Goods are batch-captured at end of shift. SAP posting timestamp records 17h00 for goods that arrived at 09h00. On-time delivery calculation is impossible without reliable delivery timestamps. Bronze layer workaround: four timestamp columns per GRN record — sap_posting_timestamp, claimed_delivery_date, grn_capture_timestamp, and data_ingested_timestamp.

Issue 2 — Supplier Contracts in PDF

Contracted rates, lead times, and payment terms exist only as unstructured PDF documents. Price variance against contracted rates cannot be calculated without structured data extraction. ADR-003 raised.

Issue 3 — Invoice Exceptions Queue

847 invoices in exceptions queue — some six months old. Managed in Excel by Finance creditors team. This is a process problem not a structural one — process improvement will resolve it. The platform makes the queue visible and measurable.

Architectural Implications

- Bronze layer: Four timestamp columns on every GRN record to separate event time from processing time.
- ADR-001: Master Product Identifier Strategy — SAP stock code must resolve to POS code and Shopify SKU. Procurement domain directly affected.
- ADR-003: Supplier Contract PDF Extraction — LLM-assisted hybrid approach recommended.
- Platform constraint: No new procurement software approved until data platform ROI demonstrated. Design must work with current broken process.

3.2 Inventory Management — Zanele Mokoena, National Inventory Manager

Interview complete — 17 April 2026 | ADRs triggered: ADR-001 | BPA-FINDING-001 raised

Process Owner Profile

Zanele Mokoena is the National Inventory Manager responsible for stock management across 87 stores and one distribution centre. She is precise and data-driven, managing inventory without reliable real-time visibility. She brought her manually maintained Monday morning replenishment spreadsheet — three hours of weekly work the platform must replace. Her domain problems are structural rather than process-driven, which makes them harder to solve and more important to formally acknowledge before go-live.

Key Business Processes

Zanele cannot reliably report how many units of any SKU exist in any store at any moment. Three data gaps undermine every inventory metric. Shrinkage is captured weekly on paper or Excel per store, emailed Monday mornings — not integrated into any replenishment system. Inter-store transfers are captured in SAP, sometimes two days after the transfer occurred. Customer returns are captured as negative sales in POS, but damaged returns going to the DC are captured separately or not at all.

The replenishment process runs from store manager identifying low stock through a replenishment order to the DC, DC picking and dispatching, and store receipt confirmation. Zanele currently manages 340 open replenishments manually in a spreadsheet rebuilt every Monday — three hours of work.

Decisions Made — Daily, Weekly, Monthly

Daily: Open replenishment orders are monitored via spreadsheet. Store manager queries about stock availability are handled. DC picking priorities are managed. Weekly: The replenishment tracking spreadsheet is rebuilt every Monday — three hours of work. Shrinkage reports are reviewed from 87 stores. Monthly: Inventory reporting to Nomvula covering stock levels, shrinkage rates, and replenishment cycle times.

Questions Zanele Cannot Currently Answer

- How many units of SKU X are in store Y right now — with confidence?
- Which stores have the biggest gap between system stock and physical stock?
- Which product categories have the highest shrinkage rates?
- How long does a replenishment order take from request to store receipt — by store, DC, and category?
- Where in the replenishment pipeline are the 340 open orders right now?

What Good Looks Like in 12 Months

Zanele's manual Monday morning spreadsheet replaced by an automated dashboard showing: open replenishment orders by store, SKU, and pipeline stage in real time; average days in replenishment pipeline by store, DC, and product category; stores with the biggest variance between system and physical stock; product categories with highest shrinkage rates; and inter-store transfer cycle times.

Data Sources and Quality Issues

SAP ECC 6.0 provides central stock management, warehouse movements, and replenishment triggers. Oracle MICROS POS provides real-time stock decrements at point of sale across 87 stores, with schema variation per store requiring careful integration design.

Two fact tables are required to serve Zanele's domain. The Inventory Periodic Snapshot captures one row per SKU per store per day with Quantity on Hand as a semi-additive measure. The Replenishment Accumulating Snapshot captures one row per replenishment order updated at each pipeline stage — Reorder Date, DC Pick Date, Dispatch Date, and Store Receipt Date. Days in Pipeline is stored by ETL calculation including weekend exclusion logic per Kimball Chapter 6. Status is derived in a view from date columns rather than stored, to reduce inconsistency risk across 87 stores.

BPA-FINDING-001 — Inventory Domain Structural Data Integrity Risk

The Inventory Management domain data quality problems are structural — not process-driven. The replenishment trigger, shrinkage capture, and returns tracking are all broken at the operational source level. The Gold layer will produce inventory metrics that are directionally correct but quantitatively approximate until three upstream operational gaps are resolved: real-time stock position visibility at store level; integration of shrinkage data into a

governed system; and unified returns capture across POS and DC. This finding must be formally acknowledged by the business before platform go-live.

Architectural Implications

- Two fact tables required: Inventory Periodic Snapshot and Replenishment Accumulating Snapshot.
- Date dimension requires a special To Be Determined row (surrogate key -1) for unknown pipeline stage dates.
- ADR-001 dependency: Product dimension surrogate key must resolve across SAP, POS, and Shopify.
- Conformed Store dimension shared across Sales, Inventory Periodic Snapshot, and Replenishment Accumulating Snapshot.

3.3 Sales and Marketing — Rethabile Sithole, Chief Commercial Officer

Interview complete — 17 April 2026 | ADRs triggered: ADR-002 (expanded), ADR-006 | BPA-FINDING-002 raised | Marketing domain added to BDR-001 v1.2

Process Owner Profile

Rethabile Sithole is the Chief Commercial Officer — the most senior stakeholder in the interview programme. She owns both the Sales domain and the newly created Marketing domain following BDR-001 v1.2. She arrived with exactly 30 minutes and low expectations having been through two failed vendor data platform promises. She left after 34 minutes more engaged than when she arrived. She speaks in commercial language — revenue, lifetime value, attribution, ROI. She does not care about data architecture. She cares about decisions.

Key Business Processes

RetailSphere operates in-store sales across 87 stores via Oracle MICROS POS, online sales via Shopify Plus, and campaign execution across Salesforce Marketing Cloud, Meta Business Manager, and Google Analytics 4. Customer acquisition converts new and lapsed customers through paid and owned media. Revenue is reported to the board by total, store, category, channel, and month.

Decisions Made — Daily, Weekly, Monthly

Daily: Revenue performance versus target is monitored by channel and store. Campaign performance is reviewed to identify which campaigns are driving traffic and conversion. Weekly: Channel mix is analysed — in-store versus online revenue split. Marketing spend efficiency is reviewed by cost per acquisition per channel. Monthly: The board presentation is produced covering total revenue, channel breakdown, and customer metrics. Marketing ROI is reported against R8.3 million annual spend — however each system applies its own attribution model resulting in likely double and triple counting.

Questions Rethabile Cannot Currently Answer

- Is the customer who spent R12,000 on Shopify last year the same person who shops in our stores — are they a loyalty member?
- What is our true customer retention rate across the business — not per channel?
- Which of our R8.3 million in marketing spend is driving revenue — and through which channel?

What Good Looks Like in 12 Months

Three Gold layer requirements stated directly by Rethabile: Customer Lifetime Value across all channels — not per channel. True retention and churn rate with a single agreed definition — six months no engagement signal across

purchase, browse, email open, or store visit. Closed-loop marketing attribution connecting campaign spend to revenue regardless of conversion channel.

A four-segment customer framework is required with monthly refresh: anonymous in-store customers representing 66% of store transactions; online-only buyers; loyalty-only members; and full cross-channel customers who are the highest value segment. SCD Type 2 is required on the Customer dimension segment attribute to track movement between segments and the triggers that caused each move.

Data Sources and Quality Issues

Sales domain sources: Oracle MICROS POS for in-store transactions, Shopify Plus for online orders, and Salesforce CRM for loyalty-linked sales. Three new source systems were identified during this interview that were absent from BDR-001 v1.0: Salesforce Marketing Cloud for campaign sends and engagement, Meta Business Manager for paid social spend and attribution, and Google Analytics 4 for online behavioural data and conversion tracking.

Issue 1 — No Cross-Channel Customer View

Salesforce Loyalty ID does not equal Shopify Customer ID does not equal POS receipt number. 66% of in-store transactions are completely anonymous. Identity resolution required — ADR-002.

BPA-FINDING-002 — Marketing Attribution Has No Source of Truth

Attribution data does not exist in any RetailSphere source system. The data platform will create it for the first time. RetailSphere will spend R8.3 million annually based on numbers that have no external validation source. ADR-006 must be formally agreed and signed by the CCO before any Gold layer marketing build begins.

Architectural Implications

- Marketing added as sixth domain in BDR-001 v1.2 — Rethabile owns acquisition, Priya Naidoo owns retention.
- ADR-002 scope expanded: Customer identity resolution must include segment classification and anonymous customer handling.
- ADR-006 raised: Marketing Attribution Logic — attribution window, model type, and cross-channel weighting must be agreed before Gold layer build.
- Customer dimension: SCD Type 2 required on segment attribute with full segment history.
- Churn definition agreed: six months no engagement signal across all channels.

3.4 Customer Loyalty — Priya Naidoo, Head of CRM and Loyalty

Interview complete — 17 April 2026 | ADRs triggered: ADR-002, ADR-007 | BPA-FINDING-003 raised

Process Owner Profile

Priya Naidoo is the Head of CRM and Customer Loyalty. She owns the CRM and Customer Loyalty Programme that has 2.1 million registered members, focusing on customer retention exclusively. She is also the Data Steward of the Customer dimension, having authority over customer identity, attributes, and business rules. Following further discussions during Step 2 stakeholder interviews, a Marketing domain was established and campaign execution and customer acquisition were moved to that domain, led by Rethabile Sithole. This boundary clarifies accountability and responsibility, with Priya's domain responsibility limited to customer retention.

Key Business Processes

After a customer is registered as a loyalty member their information — ID number, address, purchase behaviour, and communication preference — is loaded onto Salesforce CRM. Currently Priya monitors member engagement reactively

through system-generated reports that only show problems after they have occurred, as evidenced by 14,847 members unsubscribed in Q1 2026 citing excessive communication frequency. The data platform will flag at-risk members before they leave through the Contact History fact table linked to last purchase date, last category bought from, and last engagement date. There is also a POPIA compliance requirement that when a member unsubscribes, they must be recorded and excluded within four hours from all campaigns across all channels.

Decisions Made — Daily, Weekly, Monthly

Daily: Currently Priya sees unsubscribed members after they happen. The data platform will flag at-risk members before they leave. Weekly: There are currently manual campaign reports pulled from three separate systems. There will be an automated performance dashboard with all three systems in one view. Monthly: The total loyal member count is currently a reactive number. With the data platform the dashboard will show net member movement plus predicted churn for the next 30 days.

Questions Priya Cannot Currently Answer

Priya cannot currently determine the main reasons members unsubscribed before intervention could occur. This is because there is no clear picture on the 67% of members who reported too many emails and 24% who reported not relevant, in terms of how many emails were sent per member and which campaigns were not relevant.

Priya cannot currently determine whether a loyalty member is valuable enough to warrant a personalised retention intervention. This is because there is no visibility into when they last purchased, from which product category, or the last time they engaged with campaigns.

Priya cannot currently determine if the same member unsubscribed across all channels or only one. This is because there is no single source of truth or unified customer identifier.

What Good Looks Like in 12 Months

In 12 months Priya's team will be able to see every communication sent to every member in one place. When a member says stop, it stops everywhere, immediately, with no exceptions. Priya's team will be able to intervene before members leave, not after. No member will receive more than four contacts a month as Priya will know on the same day if that rule is breached and she can action immediately.

Data Sources and Quality Issues

Priya's primary source system is Salesforce CRM which holds the full loyalty member database of 2.1 million registered members including identity attributes, rewards history, and communication preferences. The Customer dimension is further enriched by three systems from the Marketing domain — Salesforce Marketing Cloud, Meta Business Manager, and Google Analytics 4 — which generate campaign engagement data that flows back into the customer profile.

Issue 1 — No Unified Customer Identifier

Priya cannot determine if the same member unsubscribed across all channels or only one. No single source of truth links Salesforce loyalty ID to Shopify customer ID or POS receipt number. ADR-002 required.

Issue 2 — No Cross-System Opt-out Communication

When a member opts out in one system, there is no mechanism to communicate that event to the other two systems. Members who have legally withdrawn marketing consent continue to receive communications through channels that have not received the opt-out signal. This is a direct POPIA compliance risk.

BPA-FINDING-003 — Domain Ownership Gap: Marketing vs Customer Loyalty

Marketing domain boundary formalised in BDR-001 v1.2. Rethabile owns acquisition. Priya owns retention. Risk: without ongoing clarity, both stakeholders may attempt to define conflicting requirements for Salesforce Marketing Cloud data in Step 3. Must be resolved before functional requirements are finalised.

Architectural Implications

Priya's interview introduced four architectural requirements. First, a Contact History fact table — one row per communication event per member per channel, sourced from all three marketing systems, enabling rolling 30-day contact frequency calculation and breach detection. Second, a tiered ingestion pipeline for POPIA compliance — opt-out events must propagate across all four systems within four hours. Third, two new dimension considerations — a Customer Preference dimension and a Campaign Frequency dimension, both governed by Priya as Data Steward. Fourth, ADR-007 — Contact Governance and Opt-out Propagation Architecture — must be formally agreed before any contact governance logic is built.

At-risk member flag requires three components: last purchase date from the Sales fact table; last category purchased derived from the Sales fact table joined to the Product dimension; and last engagement date from the Contact History fact table.

3.5 Finance — Johan van der Merwe, Group Financial Controller

Interview complete — 17 April 2026 | Key condition: 2 consecutive months parallel running before Excel decommission

Process Owner Profile

Johan van der Merwe is the Group Financial Controller with 22 years of experience. He owns the Finance domain at RetailSphere and carries personal liability for the accuracy of financial statements — incorrect reports could cost him his professional registration, reputation, and potentially carry legal consequences. He has outlasted four data platform projects at four different companies. His condition for cooperation is clear: the platform and Excel workbooks must run simultaneously for a minimum of two consecutive months before Excel is decommissioned, with all months reconciling to the cent and Johan's sign-off required.

Key Business Processes

Johan manages five core Finance processes. The month-end close uses 140 Excel workbooks built over 11 years, currently taking three days — the platform must deliver the equivalent output by 7am on day two of the new month. Three-way invoice matching in SAP against PO, GRN, and invoice currently produces 847 exceptions managed in Excel — unreconciled creditors at year-end constitute an audit finding. Late arriving invoices require a restatement process — SAP closes at midnight on the last day of the month but supplier invoices continue arriving for up to five days after month-end. Johan's team currently spends two days hunting through SAP to explain the resulting balance differences. SAP backdated postings allow transactions to carry past dates — platform extraction runs nightly and misses postings made the following day. General Ledger reconciliation is the foundational process underpinning all financial reporting.

Decisions Made — Daily, Weekly, Monthly

Daily: Johan monitors SAP for posting errors and system issues, reviews the exceptions queue for newly failed invoice matches, and checks for backdated postings that may affect prior period balances. Weekly: Creditors aging is reviewed to assess how long invoices have been in the exceptions queue, and cash flow position is monitored against payments due. Monthly: The month-end close is completed over three days, financial statements are signed off under personal liability, the procurement savings number is produced for Siphos board report, and the creditors aging analysis is produced for audit trail purposes.

Questions Johan Cannot Currently Answer

- Why does the March creditors balance on day one differ from day six — and can every rand of the difference be explained in seconds rather than two days?
- Which of the 847 exceptions queue invoices will never match and need to be written off — and how old are they?
- When SAP was extracted last night — what backdated postings were made today that the Gold layer will miss?
- Can an auditor be shown exactly where every number in the board report came from — column by column, transformation by transformation — without a week in spreadsheets?

What Good Looks Like in 12 Months

Johan will be able to check the variance between SAP and the platform in seconds rather than days through the General Ledger reconciliation report. An exportable audit trail report will be available to auditors within 24 hours of any request. When a backdated posting occurs after extraction, Johan will receive an alert within one hour containing full details — allowing him to approve or reject the restatement before the Gold layer is updated. The three-way match and the exception queue will be visible and managed within the platform, replacing the current Excel-based creditors aging process. The month-end close pack will be automated in the same format as the current Excel workbooks, delivered by 7am on day two after month-end — giving Johan and the team three days back to focus on high priority work they have been putting on hold.

Johan's instruction — recorded verbatim:

"Make sure the audit trail is section one. Not section five. Section one."

— Johan van der Merwe, Group Financial Controller

Data Sources and Quality Issues

There are three source systems in the Finance domain. SAP ECC 6.0 is the primary source holding the General Ledger, accounts payable, cost centres, accounts receivable, and all payment runs. The 140 Excel workbooks store accumulated business rules that have existed for approximately 11 years. The SQL Server Legacy DW holds two undocumented reports built in 2019, known only to Johan, that are used monthly.

The Excel workbook business rules are undocumented — a process problem. With the introduction of the platform all rules will be documented and transformed through dbt transformation models that will include version control, preventing any silent changes and providing a complete audit trail of all logic used to calculate any formula.

The Legacy DW has two undocumented reports that need to be discovered, replicated, tested, and signed off by Johan after two consecutive months to prove that the numbers are identical every month before the Legacy DW is decommissioned.

SAP allows backdated posted transactions with previous dates. The platform will extract data nightly and if there are any backdated postings Johan and the team will receive a detailed alert within one hour, allowing Johan to review and approve the restatement workflow without logging into the platform.

Architectural Implications

The reconciliation framework enforcing data lineage is the single most important architectural requirement for the Finance domain. It ensures every Gold layer number is traceable back to SAP to the cent. If any difference is detected between Bronze and SAP, the pipeline stops and an alert fires before any wrong number reaches the Gold layer. This framework is what builds Johan's trust — and Johan's trust is what makes the board trust every other number the platform produces.

The Excel business rules will be hosted within dbt transformation models. All logic and formulas will be visible, version controlled, with no silent changes. This provides Johan and external auditors with a permanent, auditable evidence trail for every calculation.

The transaction date and posting date are the two most important dates in every financial record. When a transaction occurred and when it was posted in SAP — capturing both gives the platform the ability to correct and restate invoice periods without misaligning totals.

The Legacy DW decommission follows a seven-stage migration plan: Stage 1 — Discovery with Johan's team; Stage 2 — Source data replication; Stage 3 — Report replication; Stage 4 — Testing for identical numbers; Stage 5 — Two consecutive months parallel running with both reports active simultaneously; Stage 6 — Johan's sign-off on each month; Stage 7 — Decommission of the Legacy DW.

4. Cross-domain Findings

4.1 Structural versus Process Problems

A critical insight from Step 2 is the distinction between process problems and structural problems across domains. Process problems are caused by people not following correct steps — they are fixable through discipline and process improvement. Structural problems are broken by design and require architectural solutions or operational change beyond the data platform scope.

Domain	Process Problems — Fixable	Structural Problems — Architectural Solution Required
Procurement	GRN batch capture — training can improve. Invoice exceptions — process improvement can clear.	Supplier contracts in PDFs — architectural extraction solution required (ADR-003).
Inventory	Inter-store transfer lag — process improvement can reduce.	Replenishment trigger unreliable. Shrinkage in Excel. Returns split across two systems. All structural — BPA-FINDING-001.
Sales	66% anonymous transactions — loyalty card adoption is a process improvement.	No unified customer ID across systems — structural. Requires ADR-002.
Marketing	Attribution methodology alignment — process agreement.	No attribution source of truth exists anywhere — structural. Platform creates it first time. ADR-006.
Finance	Business rules trapped in Excel — documentation process can fix.	SAP backdated postings — SAP designed to allow it. Architectural workaround required.

5. Formal Findings and Risks

BPA-FINDING-001 — Inventory Domain Structural Data Integrity Risk

The Inventory Management domain data quality problems are structural. The replenishment trigger, shrinkage capture, and returns tracking are broken at the operational source level. Gold layer inventory metrics will be directionally correct but quantitatively approximate until the business resolves three structural gaps. Formal business acknowledgement required before platform go-live.

BPA-FINDING-002 — Marketing Attribution Has No Source of Truth

Attribution data does not exist in any RetailSphere source system. The data platform will create it for the first time. RetailSphere will make R8.3 million in annual spend decisions based on numbers with no external validation source. ADR-006 must be formally agreed and signed by the CCO before any Gold layer marketing build begins.

BPA-FINDING-003 — Domain Ownership Gap: Marketing vs Customer Loyalty

Marketing domain boundary formalised in BDR-001 v1.2. Rethabile Sithole owns acquisition. Priya Naidoo owns retention. Risk: without clear ongoing communication, both stakeholders may attempt to define conflicting requirements for Salesforce Marketing Cloud data in Step 3.

BPA-FINDING-004 — Unified Customer Identity Is the Critical Path Dependency

ADR-002 — Unified Customer Identity Resolution — is the single most important functional requirement across the platform. Every customer-level metric from every domain is blocked until it is resolved: Sipho's customer-driven procurement insights, Zanele's demand-linked replenishment, Rethabile's CLV and attribution, Priya's at-risk flag and breach detection, and Johan's revenue by customer segment. ADR-002 must be resolved in Step 6 before Step 7 dimensional modelling can be finalised.

6. ADRs Triggered by Step 2

ADR	Decision Required	Triggered By	Priority
ADR-001	Master Product Identifier Strategy	<i>BDR-001 v1.0 — no SAP/POS/Shopify product link</i>	HIGH — blocks Step 7 Product dimension
ADR-002	Unified Customer Identity Resolution	<i>BDR-001 v1.0 + Rethabile interview (scope expanded)</i>	CRITICAL — blocks Step 7 Customer dimension
ADR-003	Supplier Contract PDF Extraction	<i>Sipho Dlamini interview</i>	MEDIUM — blocks price variance metric
ADR-006	Marketing Attribution Logic	<i>Rethabile Sithole interview</i>	CRITICAL — must be agreed before Gold layer marketing build
ADR-007	Contact Governance and Opt-out Propagation	<i>Priya Naidoo interview</i>	HIGH — POPIA legal obligation

7. Next Steps

- 1. CDO review and sign-off on BPA-001 v1.0.
- 2. Formal business acknowledgement of BPA-FINDING-001 — Inventory structural data integrity risk.
- 3. Formal CCO sign-off on BPA-FINDING-002 — ADR-006 to be written before any Gold layer marketing build.
- 4. Communication and confirmation of Marketing/Loyalty domain boundary with both Priya Naidoo and Rethabile Sithole before Step 3.
- 5. Proceed to Step 3 — Functional Requirements Definition using all five interview findings as source material.
- 6. Initiate ADR-001 and ADR-002 formal write-up as priority items in Step 6 — both are critical path dependencies for Step 7 dimensional model design.

8. Document Change Log

Version	Date	Author	Description
0.1	17 Apr 2026	T. Mbalo	Initial draft — Sipho, Zanele, Rethabile interviews complete. Priya and Johan pending.
0.2	17 Apr 2026	T. Mbalo	Priya Naidoo interview added. ADR-007 triggered. BPA-FINDING-003 raised.
0.3	17 Apr 2026	T. Mbalo	Johan van der Merwe interview added. All five interviews complete.
1.0	1 May 2026	T. Mbalo	Final version. All sections complete. Four formal findings. Cross-domain analysis added. Pending CDO sign-off.

— End of Document —

BPA-001 v1.0 | RetailSphere Data Platform Programme | Confidential